

Figure 4: Preference learning convergence
KL divergence vs. L1 updates, mean $\pm 1\sigma$ over 50 runs

